




**International University for Science &  
Technology (IUST)**  
**Department of Computer & Informatics  
Engineering**  
**Neural Networks unit (7)**

# Deep Learning

Neural Networks

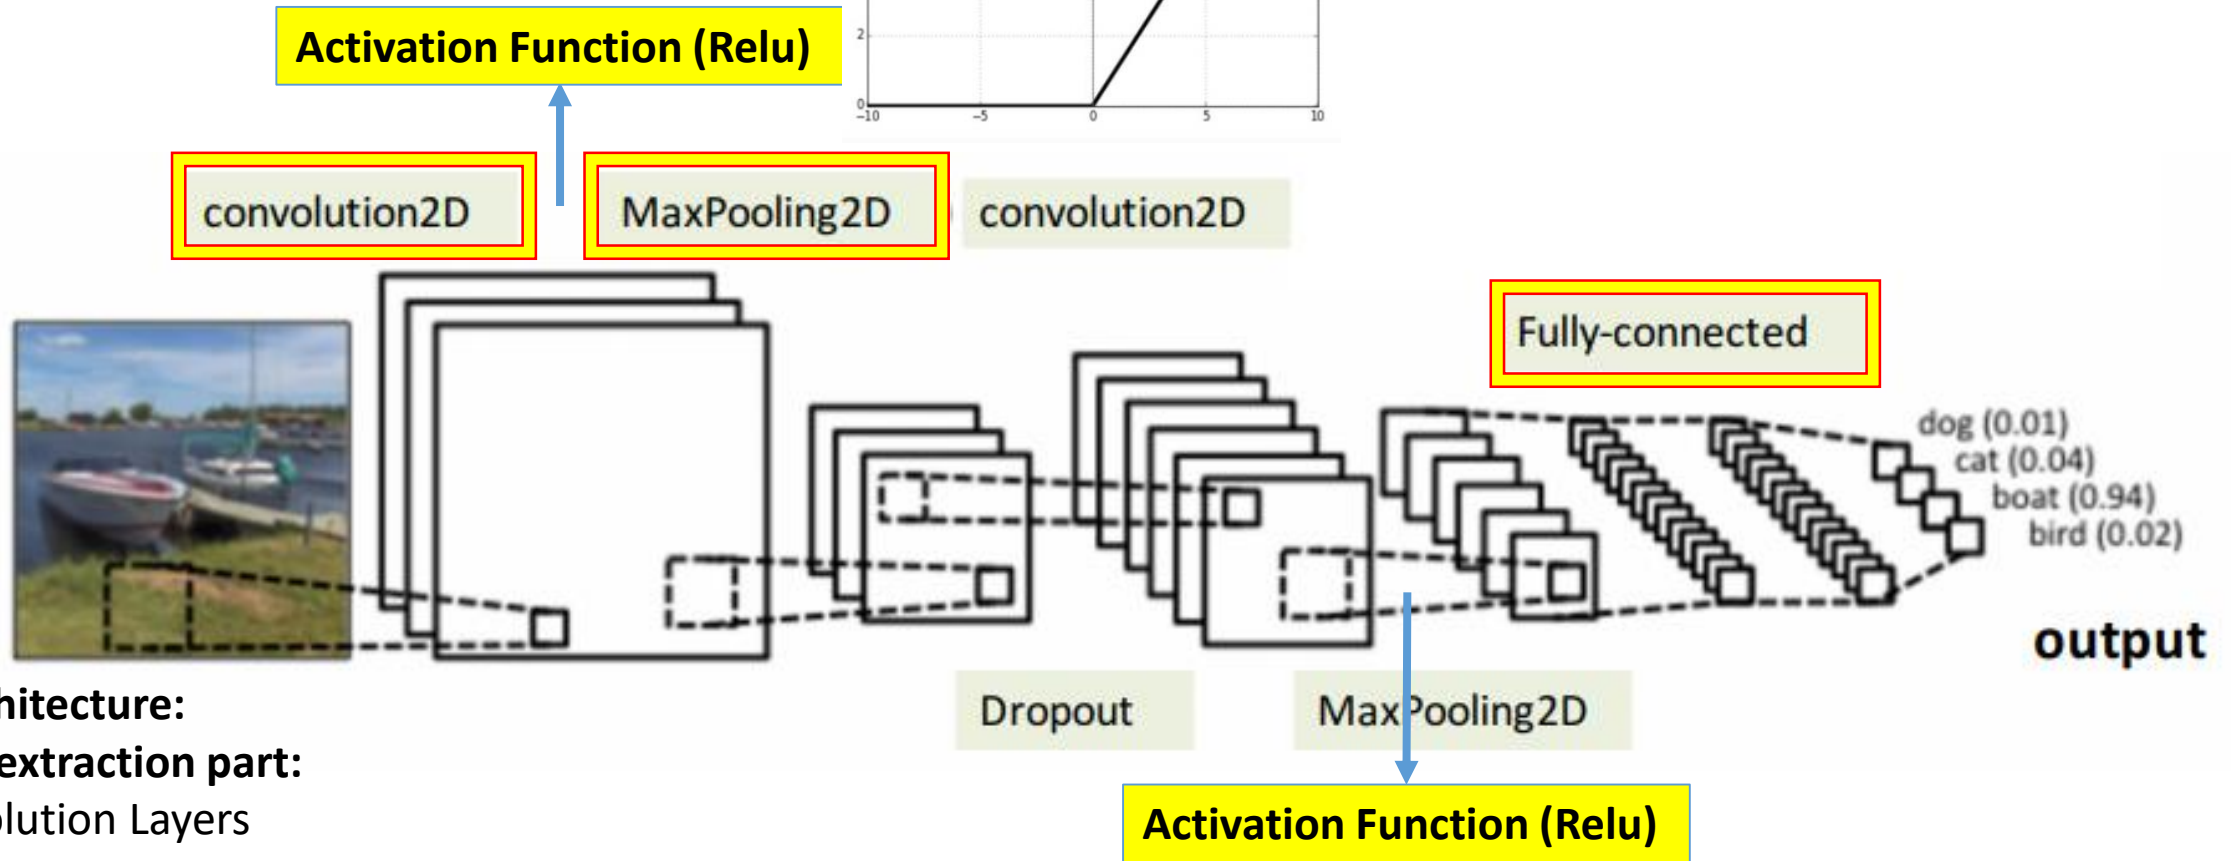
Dr. Ali Mayya

# Machine Learning Vs. Deep Learning



# Deep Learning- CNN and Recurrent Neural Networks

- 



**CNN architecture:**

**Feature extraction part:**

- 1- Convolution Layers
- 2- Pooling Layers

**Classification part:**

- 1- Fully-Connected layers:

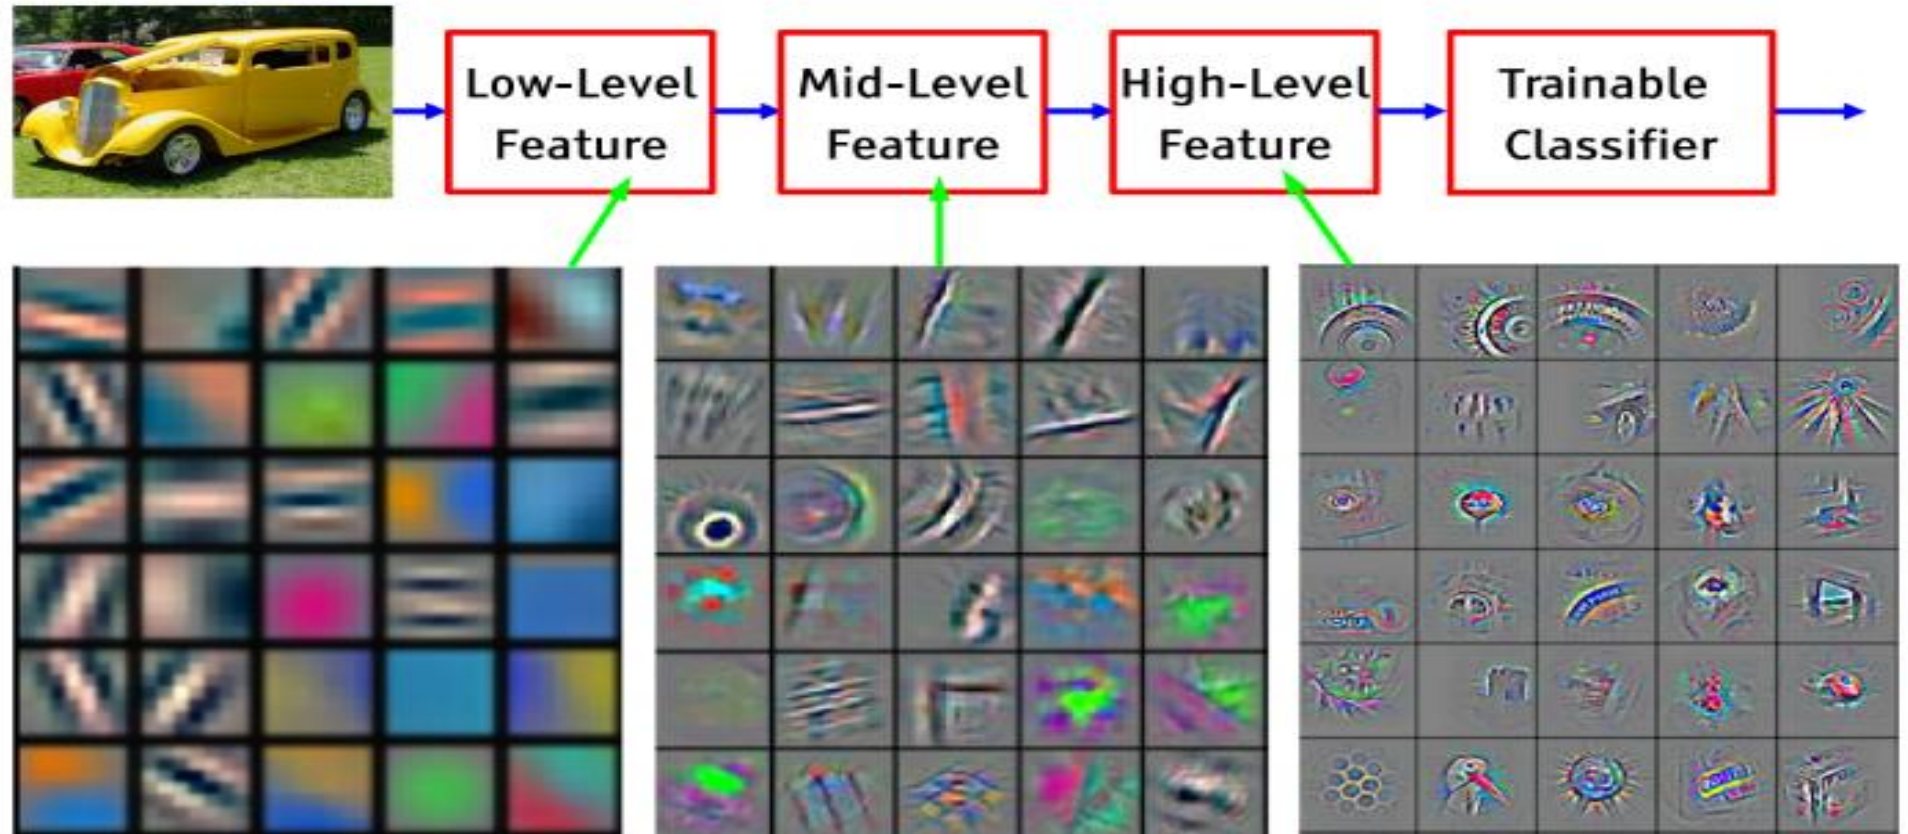
Some dense layers followed by

final layer is the classification layer with sigmoid or Softmax activation function.

$$\text{softmax}(z)_i = \frac{\exp(z_i)}{\sum_j \exp(z_j)}$$

# Deep Learning- CNN Convolutional Neural Networks

## Hierarchical representation




Feature visualization of convolutional net trained on ImageNet from [Zeiler & Fergus 2013]



# Deep Learning- CNN Convolutional Neural Networks

## Convolution



Convolution layer contains main part called filter (kernel) which is a  $k \times k$  matrix applied to the input image in a specific neighborhood

What is convolution?

The weighted sum of multiplication between kernel and corresponding neighborhood of the image's pixel

Based on the values of the kernel, the output of the convolution is defined.

# Deep Learning- CNN Convolutional Neural Networks

## Convolution

- Based on the values of the kernel, the output of the convolution is defined.

Derivation on both x and y (sum of kernel values=1) → Horizontal and vertical edge sharpening

Derivation on both x and y (sum of kernel values=0) → Horizontal and vertical edge detection → isolated points detection

Derivation on y axis (sum of values=0) and 4 neighborhood is multiplied by 2 → Strong Horizontal edges

Sum of values=1 (all positives) and apply the weighted sum → Smoothing



Original




Sharpen



Edge Detect




“Strong” Edge Detect



Smoothing


# Deep Learning- CNN Convolutional Neural Networks

## Convolution




No padding, stride 1

Input image size:  $4 \times 4$   
Kernel ( $3 \times 3$ ),  
stride=1, padding=  
No  
Output size:  $2 \times 2$



Padding 1, stride 1

Input image size:  $5 \times 5$   
Kernel ( $3 \times 3$ ),  
stride=1, padding= 1  
Output size:  $5 \times 5$



No padding, stride 2

Input image size:  $5 \times 5$   
Kernel ( $3 \times 3$ ),  
stride=2, padding=  
No  
Output size:  $2 \times 2$

Convolution has two parameters:  
**Stride:** how many squares the kernel skip when moving across image, from left to right and from top to bottom.  
**Padding:** number of added rows and columns to the borders of the image to handle the neighborhood pixels.  
The output of each convolution layer is called the activation maps

All activation maps's values are subjected to Relu activation function


# Deep Learning- CNN Convolutional Neural Networks

## Convolution

- Output volume size=  $N*((W-K+2P)/S+1)$   
N: number of filters of the convolutional layer  
W: image size  
K: kernel size  
P: padding, S: stride

### Convolution layer

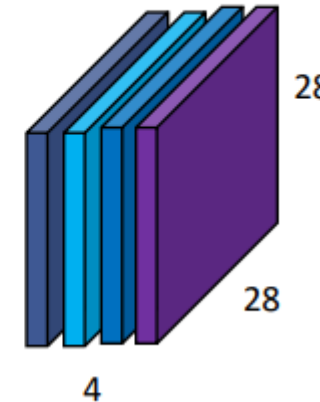
32×32×3 image



If there are **four 5×5×3 filters**

Convolve (slide) over all spatial locations

4 separate activation maps



Output volume  
size=  
 $4*((32-5+0)/1+1) =$   
4 activation maps of  
size 28\*28



# Deep Learning- CNN Convolutional Neural Networks

## Convolution

- 

$$\text{Output volume size} = N * ((W - K + 2P) / S + 1)$$

Input size: 32\*32, Kernel size: 5\*5, padding=2, stride = 1

$$\text{Output volume size} = (32 + 2 \times 2 - 5) / 1 + 1 = 32 \text{ spatially}$$

How to compute number of parameters of each layer?


For each filter  $k * k * 3 + 1$  (3: colored image, 1: Bias)

For all filters of the convolution layer: number of filters \* number of parameters per layer

In this example:

Each filter has  $5 * 5 * 3 + 1$  (for bias) = 76 parameters

For all 10 filters  $\rightarrow$  760 parameters





# Deep Learning- CNN Convolutional Neural Networks

## Convolution


•


$$\text{Output volume size} = N * ((W - K + 2P) / S + 1)$$






fewer parameters!





Fewer parameters

Even fewer parameters

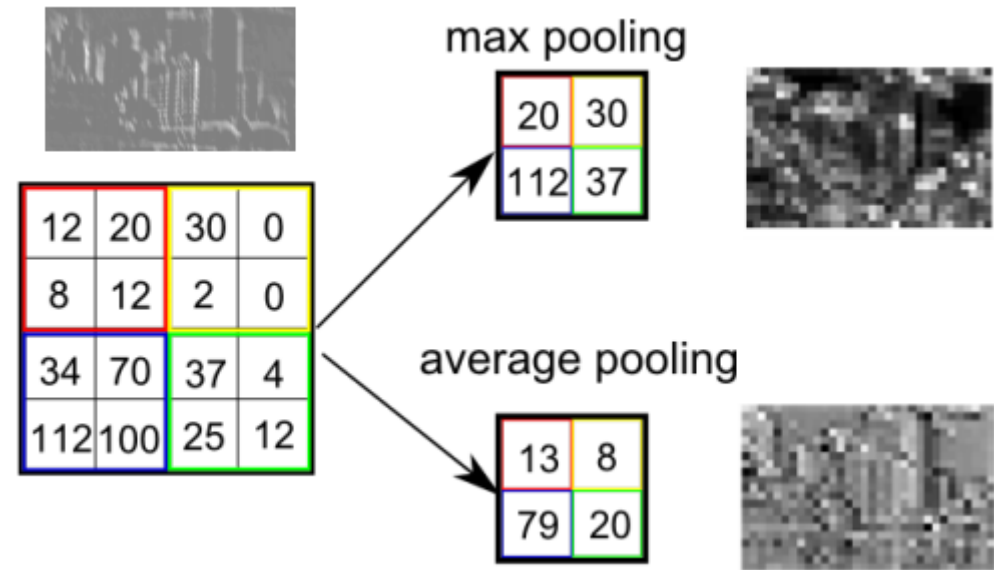




Shared weights

# Deep Learning- CNN Convolutional Neural Networks

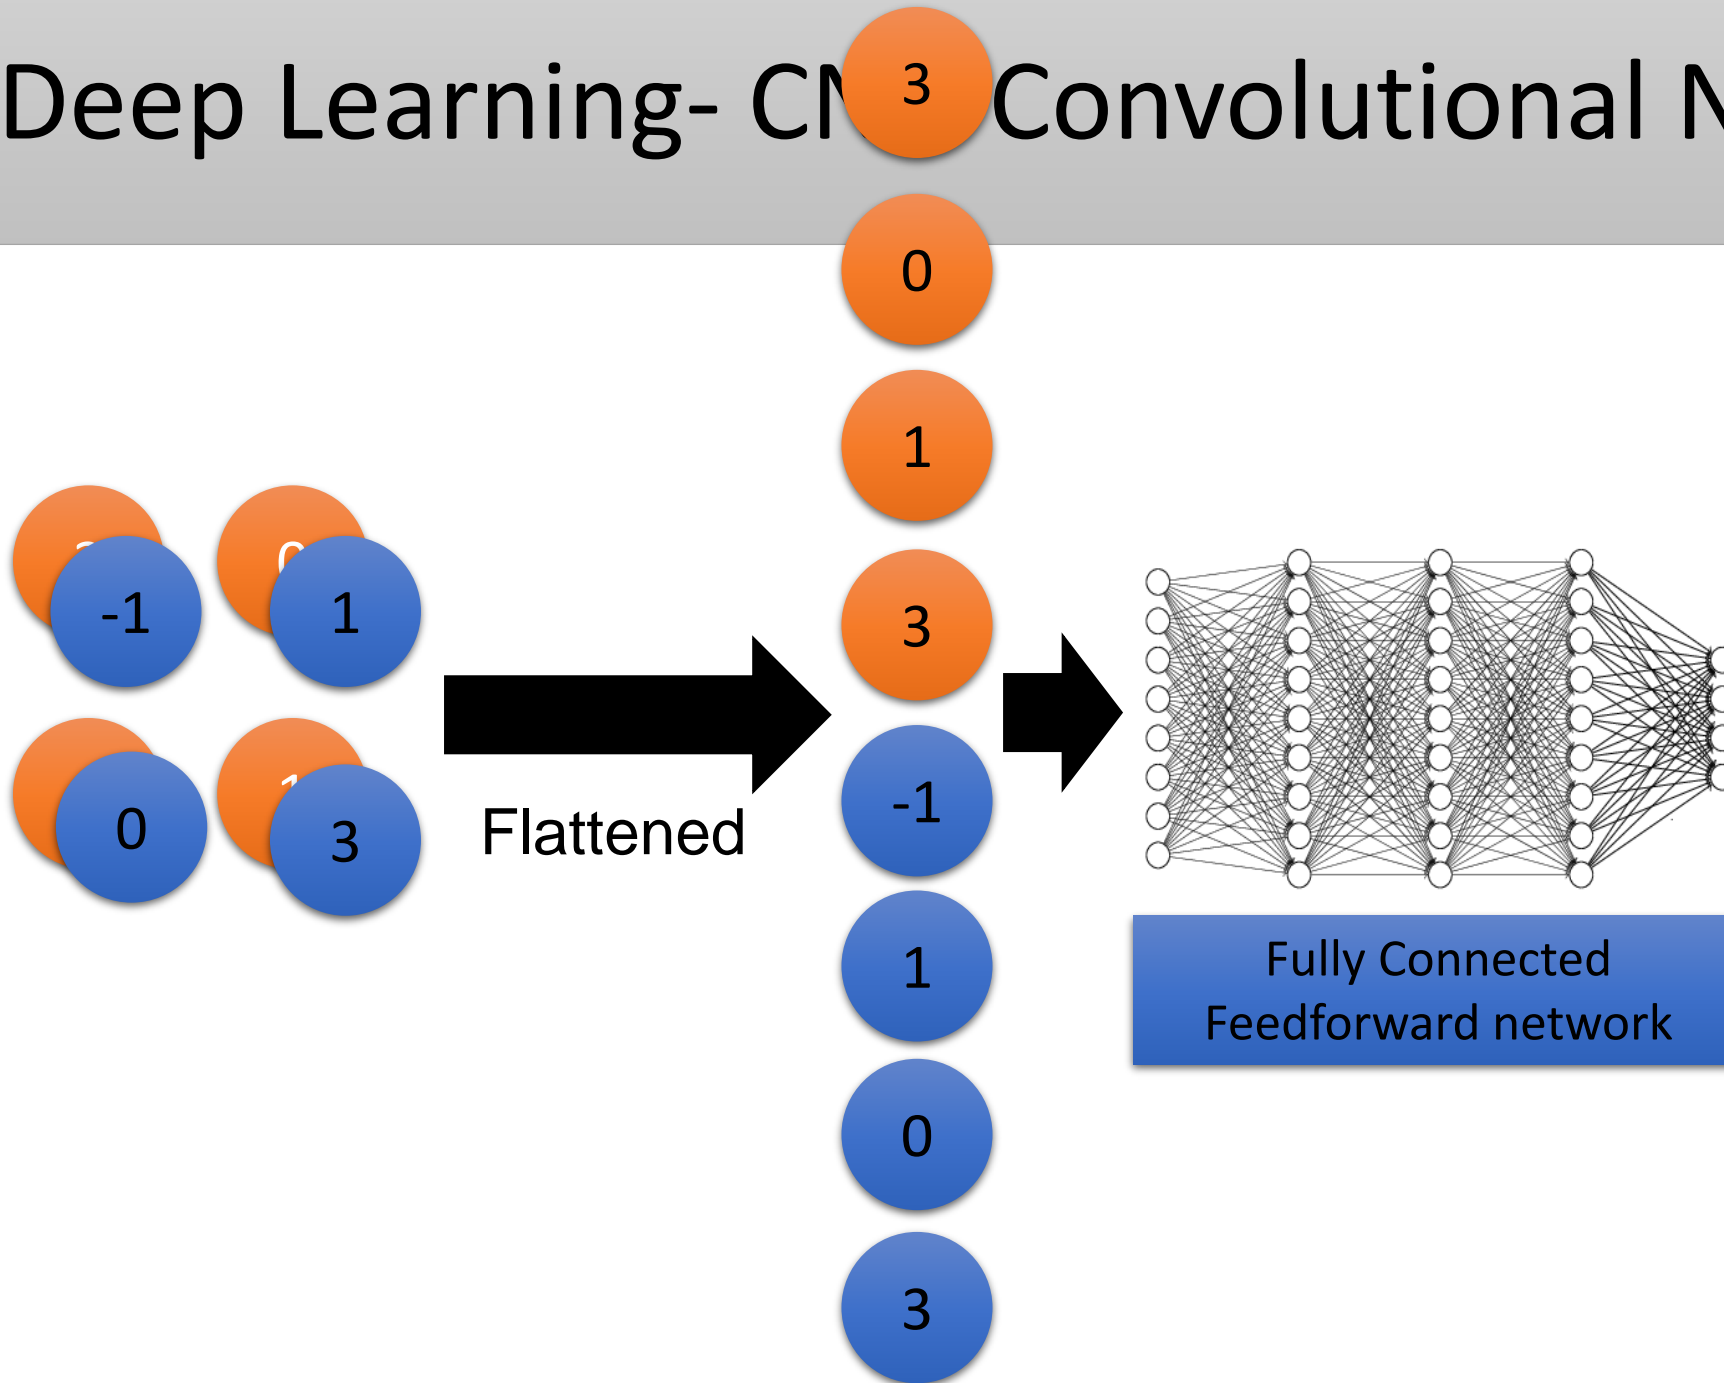
## Pooling

**Minimize the activation maps size (lower computations) and maintain the best feature information**





# Deep Learning- CNN Convolutional Neural Networks



## Flattening and FC

**Convert the last activation map of the feature extraction part into 1D vector**

**And then add fully-connected layers as needed**



**The final fully-connected layer is the output (classification layer)**

**Number of neurons in the classification layers equals to the number of required outputs.**

**Activation function: Sigmoid (binary classification)**

**Softmax (multi-class classification)**

# Deep Learning- CNN Convolutional Neural Networks



## Prevent overfitting

**We can reduce the effect of overfitting by adding nonlinearity to the architecture of the deep network.**

**In the classification part, we can add the dropout layers**

**Dropout layers can be inserted between two adjacent dense layers to remove some neurons with a specific rate (for example 25%)**

**Dropout 25% means dropping out 25% of the neurons of the current layer randomly (set their values to 0).**

# The whole CNN- Cat dog classification example



Convolution

Max Pooling

Convolution


Max Pooling

Can  
repeat  
many  
times

Flattened

cat dog .....

Fully Connected  
Feedforward network





# Deep Learning- CNN applications

- 

## Image Captioning

No errors

Minor errors

Somewhat related



*A white teddy bear sitting in the grass*



*A man in a baseball uniform throwing a ball*

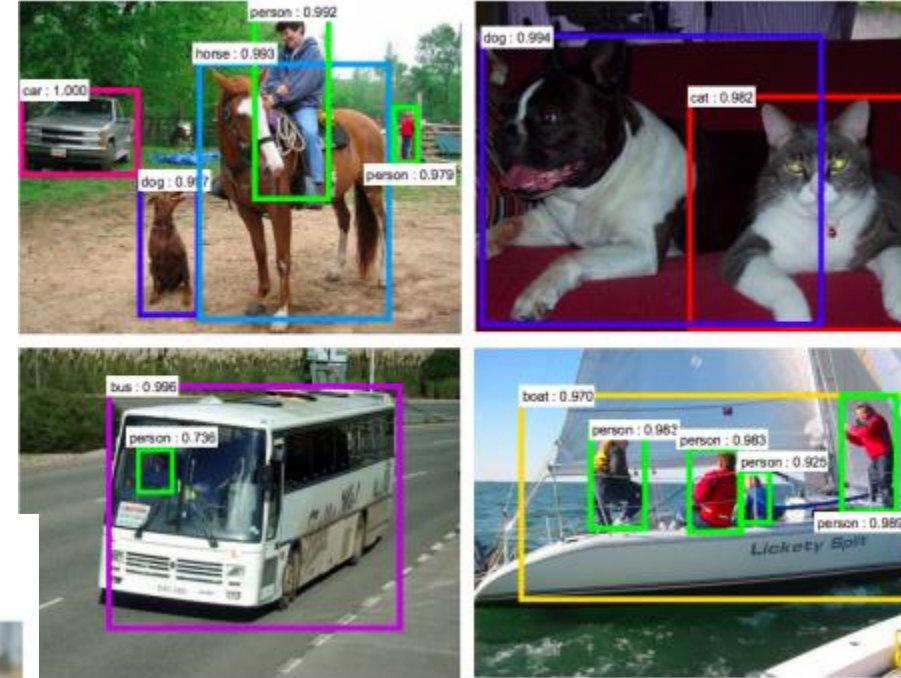


*A woman is holding a cat in her hand*

## Self-driving cars



## Detection [Ren et al., 2015]



## Image Recognition

# Deep Learning- Why GPU?

- GPU has the power of parallel processing allowing to handle batch of images at the same time reducing the training time

```
X_train shape: (50000, 3, 32, 32)
50000 train samples
10000 test samples
Using real-time data augmentation.
Epoch 1/200
50000/50000 [=====] 734s
Epoch 2/200
50000/50000 [=====] 733s
Epoch 3/200
50000/50000 [=====] 733s
Epoch 4/200
50000/50000 [=====] 733s
```

Running time **without** GPU

**VS**

```
X_train shape: (50000, 3, 32, 32)
50000 train samples
10000 test samples
Using real-time data augmentation.
Epoch 1/200
50000/50000 [=====] 27s
Epoch 2/200
50000/50000 [=====] 27s
Epoch 3/200
50000/50000 [=====] 27s
Epoch 4/200
50000/50000 [=====] 27s
```


Running time **with** GPU



# Some CNN famous examples

## VGG16 model

5 Groups of Conv+Maxpool layers  
Flattened layer (7\*7\*512,1) size layer  
Fully-connected layer 1\*1\*4096  
Fully-connected layer 1\*1\*4096  
Fully-connected layer 1\*1\*1000 (1000 classes of ImageNet dataset)



# Some CNN famous examples


## VGG16 model

### Transfer Learning:

Re-use of the VGG pretrained model (trained on ImageNet dataset) and re-train it on our specific dataset (cancer classification for example)

#### Two options:

- Maintain architecture with modification of the input data and the classification layer and retrain all layers.
- Maintain architecture with modification of the input data and the classification layer and retrain only the classification part (FC layers) (faster)



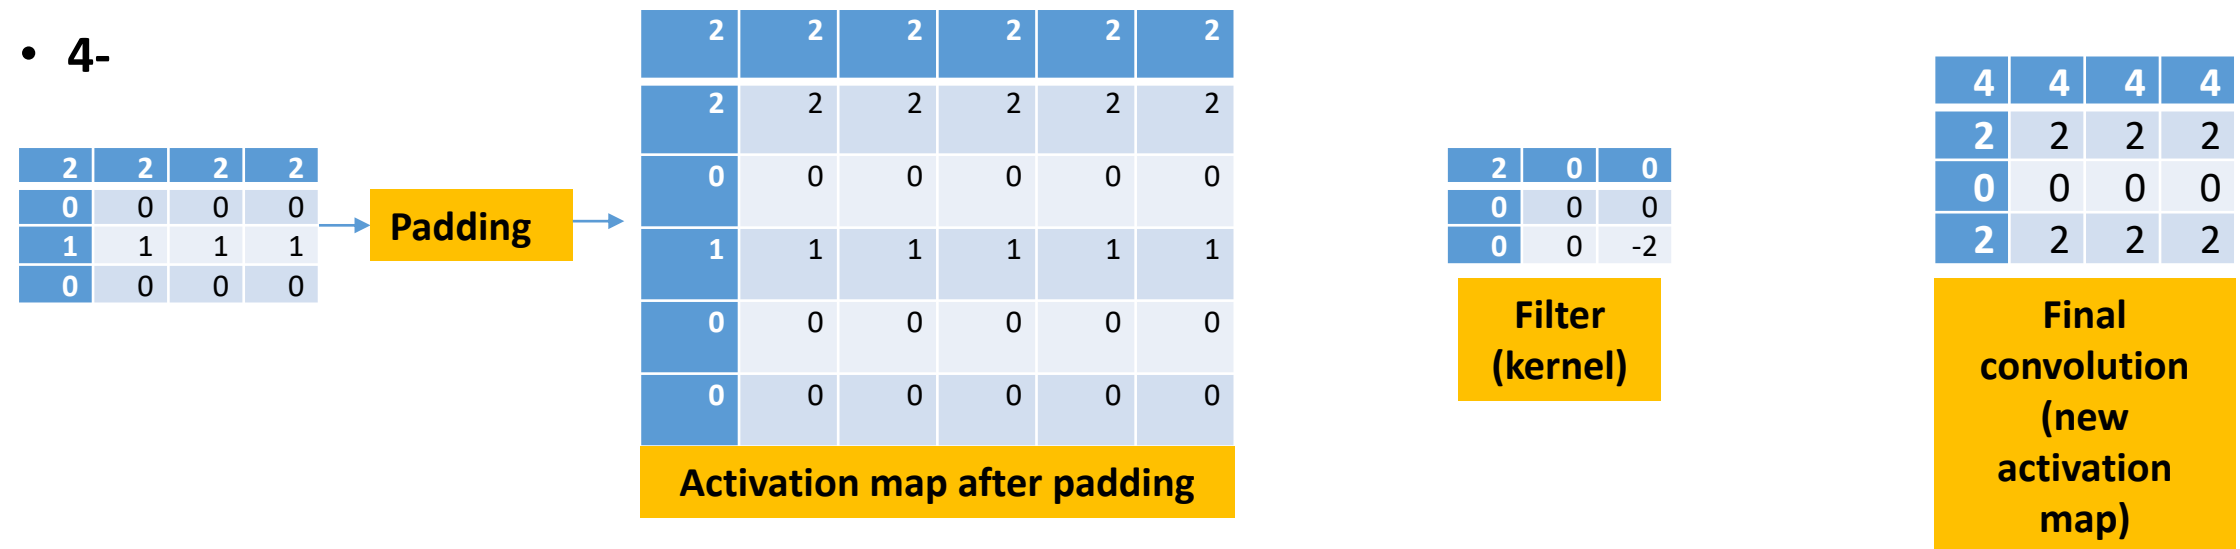
# Numerical Example

- Using a CNN with input size of  $32 \times 32$  to be classified into **four** classes (COVID-19 ,Pneumonia, Influenza, Normal) and using a **padding=1, stride=2**, two convolutional layers. Number of filters of the first layer=**8**, and for second layer=**16**. Suppose all filters of size  **$3 \times 3$** , and each conv layer is followed by a  **$2 \times 2$  max pooling** layer and has a **Relu** activation function. The classification part consists of Flatten, Dense layer of **16** neurons, and final classification dense layer. (Note: use floor for your calculations).
- What is the activation maps' size on the output of each layer (conv and max pool layers).
- What is activation function of the output layer and what are number of neurons?
- To avoid overfitting, what is the required modification of the architecture.
- Suppose we have this kernel and the activation map on the output of the second pooling layer. Apply the convolution to the image. You need to make padding by repeating rows and columns before make the convolution.
- Compute the size of the output of the final max pool, then dense layer then classification layer.

<table><tr><td>2</td><td>0</td><td>0</td></tr><tr><td>0</td><td>0</td><td>0</td></tr><tr><td>0</td><td>0</td><td>-2</td></tr></table> <p>Filter المرشح</p>	2	0	0	0	0	0	0	0	-2	<table><tr><td>2</td><td>2</td><td>2</td><td>2</td></tr><tr><td>0</td><td>0</td><td>0</td><td>0</td></tr><tr><td>1</td><td>1</td><td>1</td><td>1</td></tr><tr><td>0</td><td>0</td><td>0</td><td>0</td></tr></table> <p>Max الصورة على خرج طبقة الثانية pooling</p>	2	2	2	2	0	0	0	0	1	1	1	1	0	0	0	0	<table><tr><td>4</td><td>4</td><td>4</td><td>4</td></tr><tr><td></td><td></td><td></td><td></td></tr><tr><td></td><td></td><td></td><td></td></tr><tr><td></td><td></td><td></td><td></td></tr></table> <p>الصورة على خرج طبقة الثالثة Convolution</p>	4	4	4	4												
2	0	0																																									
0	0	0																																									
0	0	-2																																									
2	2	2	2																																								
0	0	0	0																																								
1	1	1	1																																								
0	0	0	0																																								
4	4	4	4																																								

# Numerical Example

- **1-** First layer output  $(W-K+2P)/S+1 = (32-3+2)/2+1 = 8*16*16$
- Max pooling:  $8*8*8$
- Second layer output  $(W-K+2P)/S+1 = (8-3+2)/1+1 = 16*4*4$
- Max pooling:  $16*2*2$
- Flatten  $\rightarrow 64$  (4 samples per kernel)
- Dense  $\rightarrow 16$
- **2-** Activation function is softmax, number of neurons: 4 since the problem is multi-class classification.
- **3-** Avoid overfitting by adding dropout layers between the dense layers of the classification part to add non-linearity of the model.
- **4-**



# Numerical Example

- 5- output on the max pooling

4	4	4	4
2	2	2	2
0	0	0	0
2	2	2	2

**Final  
convolution  
(new  
activation  
map)**

4	4
2	2

**Output  
of the  
third  
Max  
pool  
layer**

4	4	2	2
---	---	---	---

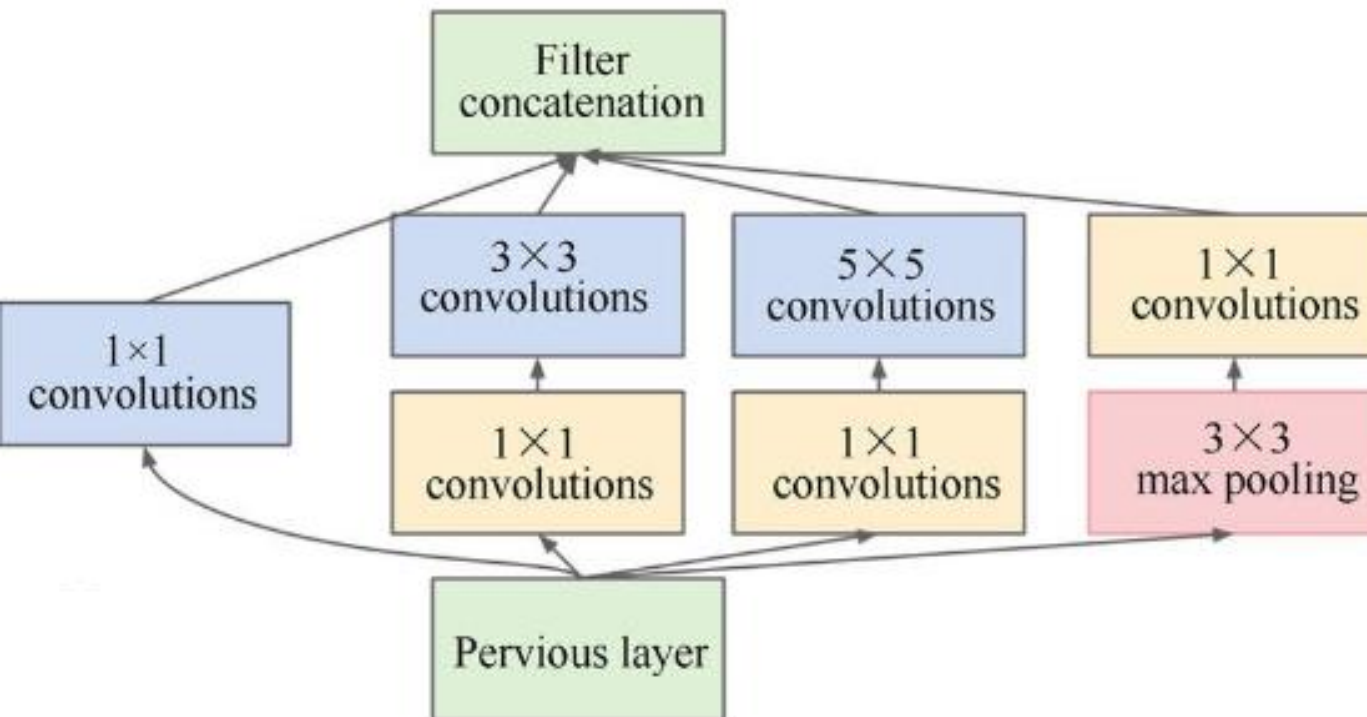
**Flatten**



# Other CNN famous examples

## Inception


Parallel architecture contains different convolution sizes ( $1 \times 1$ ,  $3 \times 3$ ,  $5 \times 5$ ) and then they are concatenated to get different scaling features of the same input image.



Inception modules, Szegedy et al

## InceptionResNet

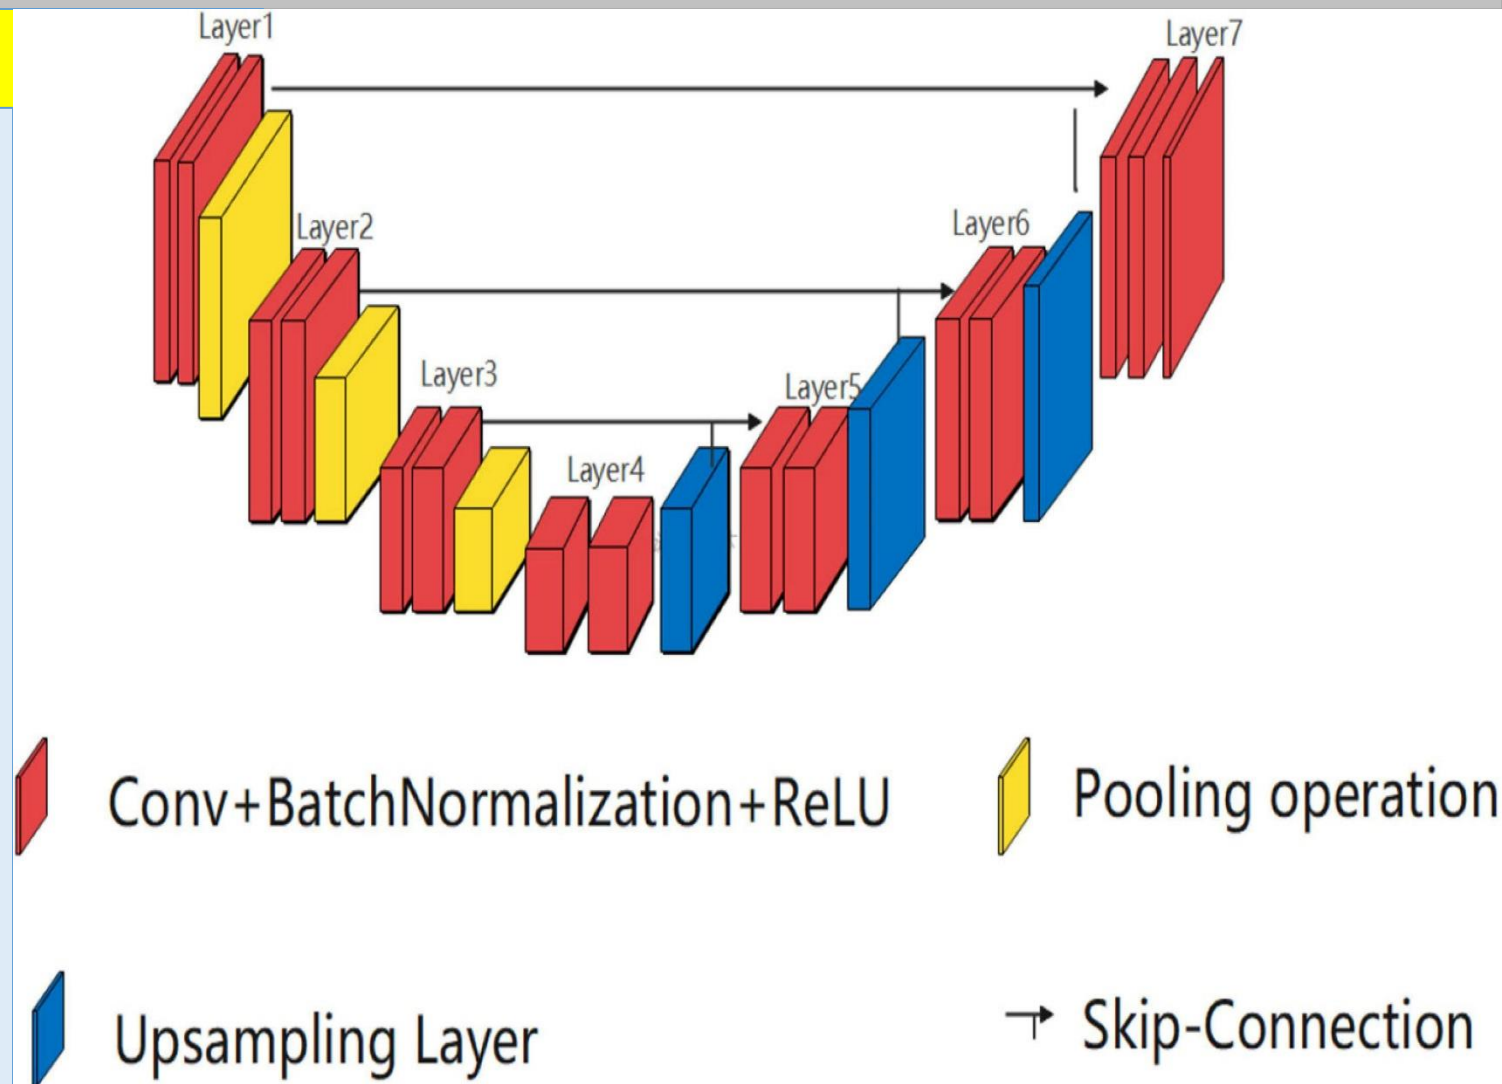
Add identity connections (called residual connections) that escape one or more convolution layers and this can allow more deep (more hidden layers) without problem of gradient vanishing.



# Other CNN famous examples

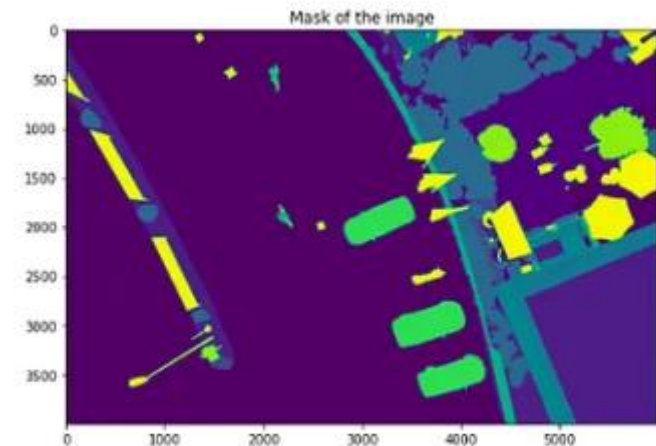

## UNet

- The U-Net architecture, introduced in **2015**, is a popular model for image segmentation tasks. It has a **U-shaped design** with two main parts:
- **Encoder (Contraction Path):**
  - This part captures context and extracts features.
  - Consists of repeated **convolutional layers** followed by **max-pooling** for downsampling.
- **Decoder (Expansion Path):**
  - This reconstructs the spatial resolution to generate the segmentation map.
  - Includes **upsampling layers** (transpose convolutions) and **skip connections** from the encoder, which provide fine-grained details.



# Other CNN famous examples

## Better illustration of UNet



# Some terms in DL field

## Optimizer

An **optimizer** is an algorithm or method used to adjust the model's parameters (e.g., weights and biases) to minimize the **loss function (error)** during training.

It determines how the model learns from the data by updating these parameters iteratively based on the gradients calculated during learning algorithm.

The most famous optimizer is “Adam” which uses the gradient descent algorithm

## What is learning algorithm of DL models?

Backpropagation (used the gradient descent learning)

### Binary classification

Binary Cross-Entropy:

$$\mathcal{L} = -\frac{1}{n} \sum_{i=1}^n [y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)]$$

## What is loss function?

- A **loss function** quantifies how well (or poorly) the model's predictions match the actual target values. The optimizer uses the gradients of the loss function to update the model parameters.

Cross-Entropy Loss:

### Multi-class classification

$$\mathcal{L} = -\frac{1}{n} \sum_{i=1}^n \sum_{j=1}^k y_{ij} \log(\hat{y}_{ij})$$